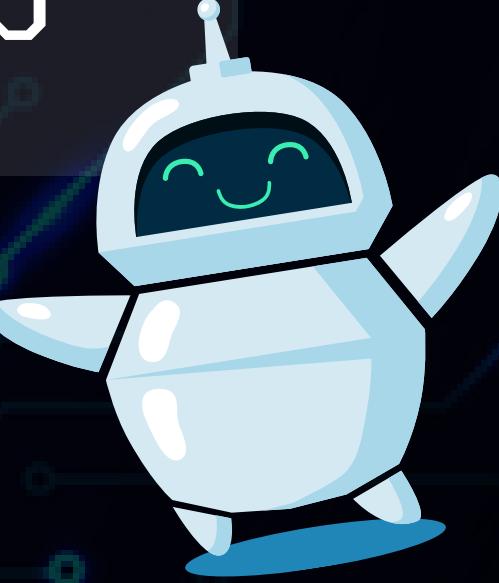




# ML ENGINE GENERATION



By: Andrew Shehata, Malak El-Khashab

# TABLE OF CONTENT

- Introduction
- Approaches
  - Our Model
    - Results, Advantages, Disadvantages
  - Pre-trained
    - Results, Advantages, Disadvantages
- Conclusion and Findings



# INTRODUCTION

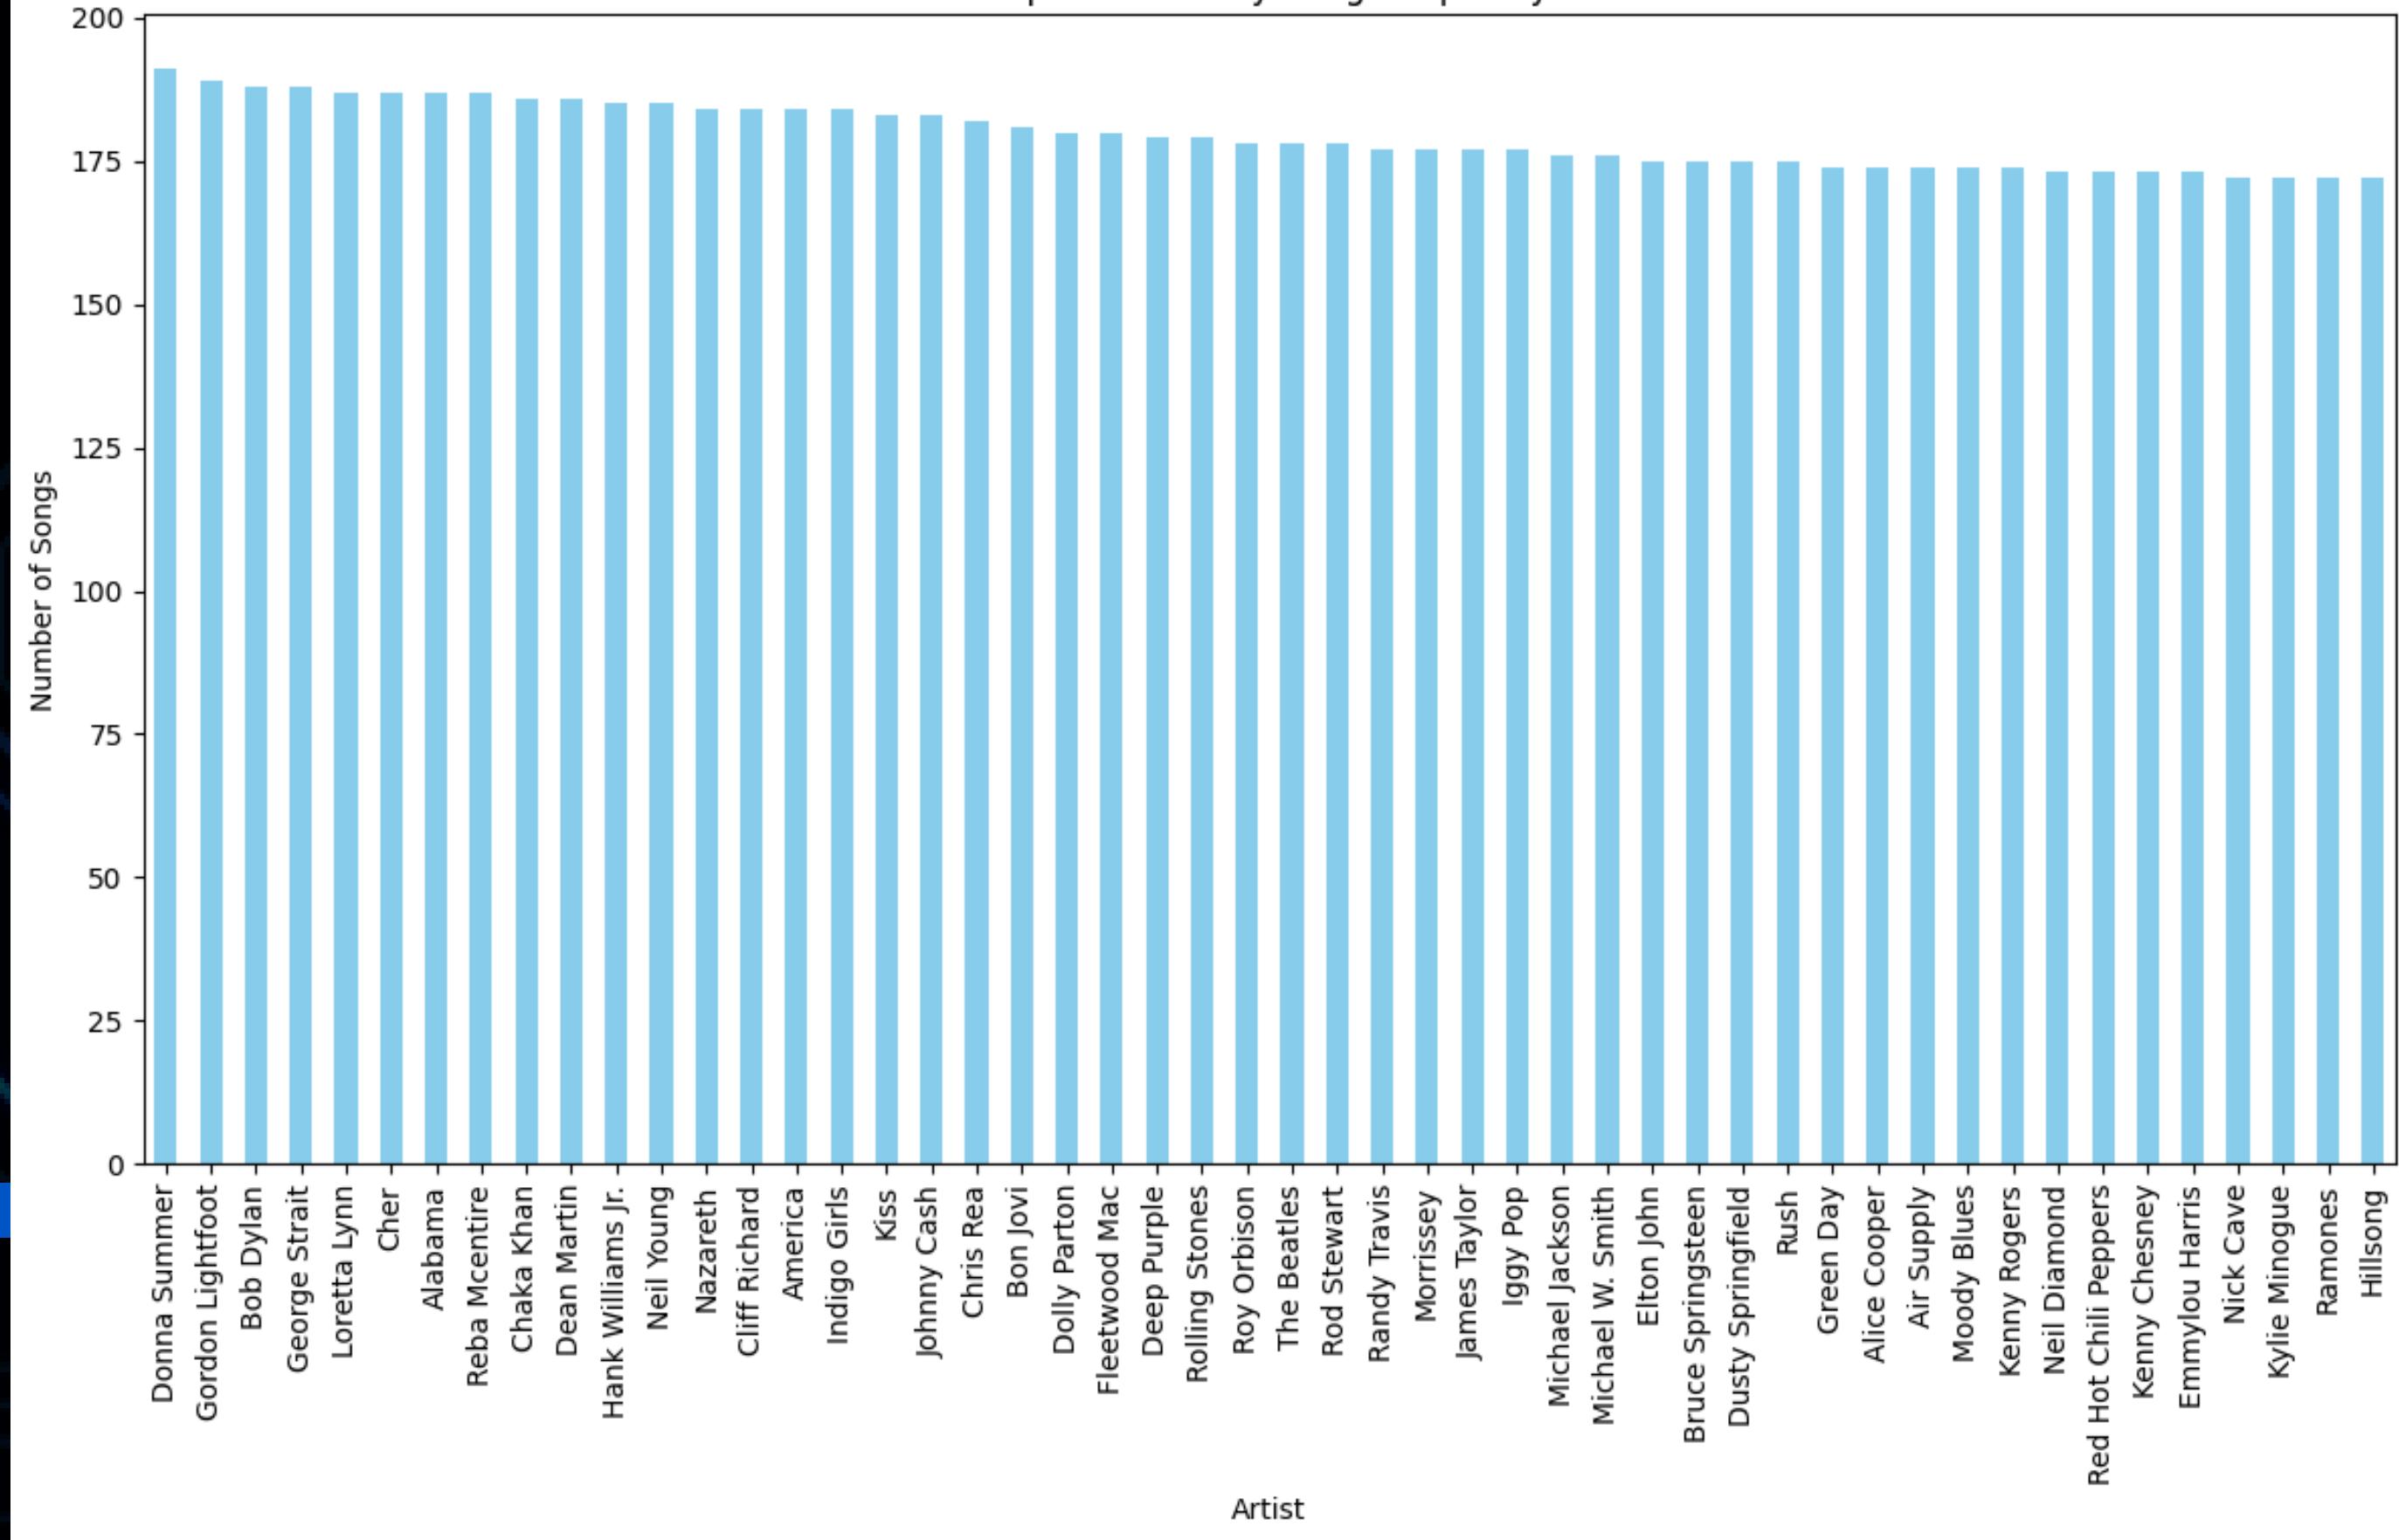
- **Project Objective:** Generate song lyrics based on a specific artist's style and given lyrics.
- **Technology Used:** Advanced NLP techniques and a pretrained T5 model.
- **Goals:**
  - Emulate the unique lyrical style of various artists.
  - Generate coherent and creative lyrics matching the input and artist's style.
- **Key Steps:**
  - Data pre-processing.
  - Model training.
  - Fine-tuning for stylistic accuracy.
- **Outcome:** Achieved relevant and stylistically accurate generated lyrics.

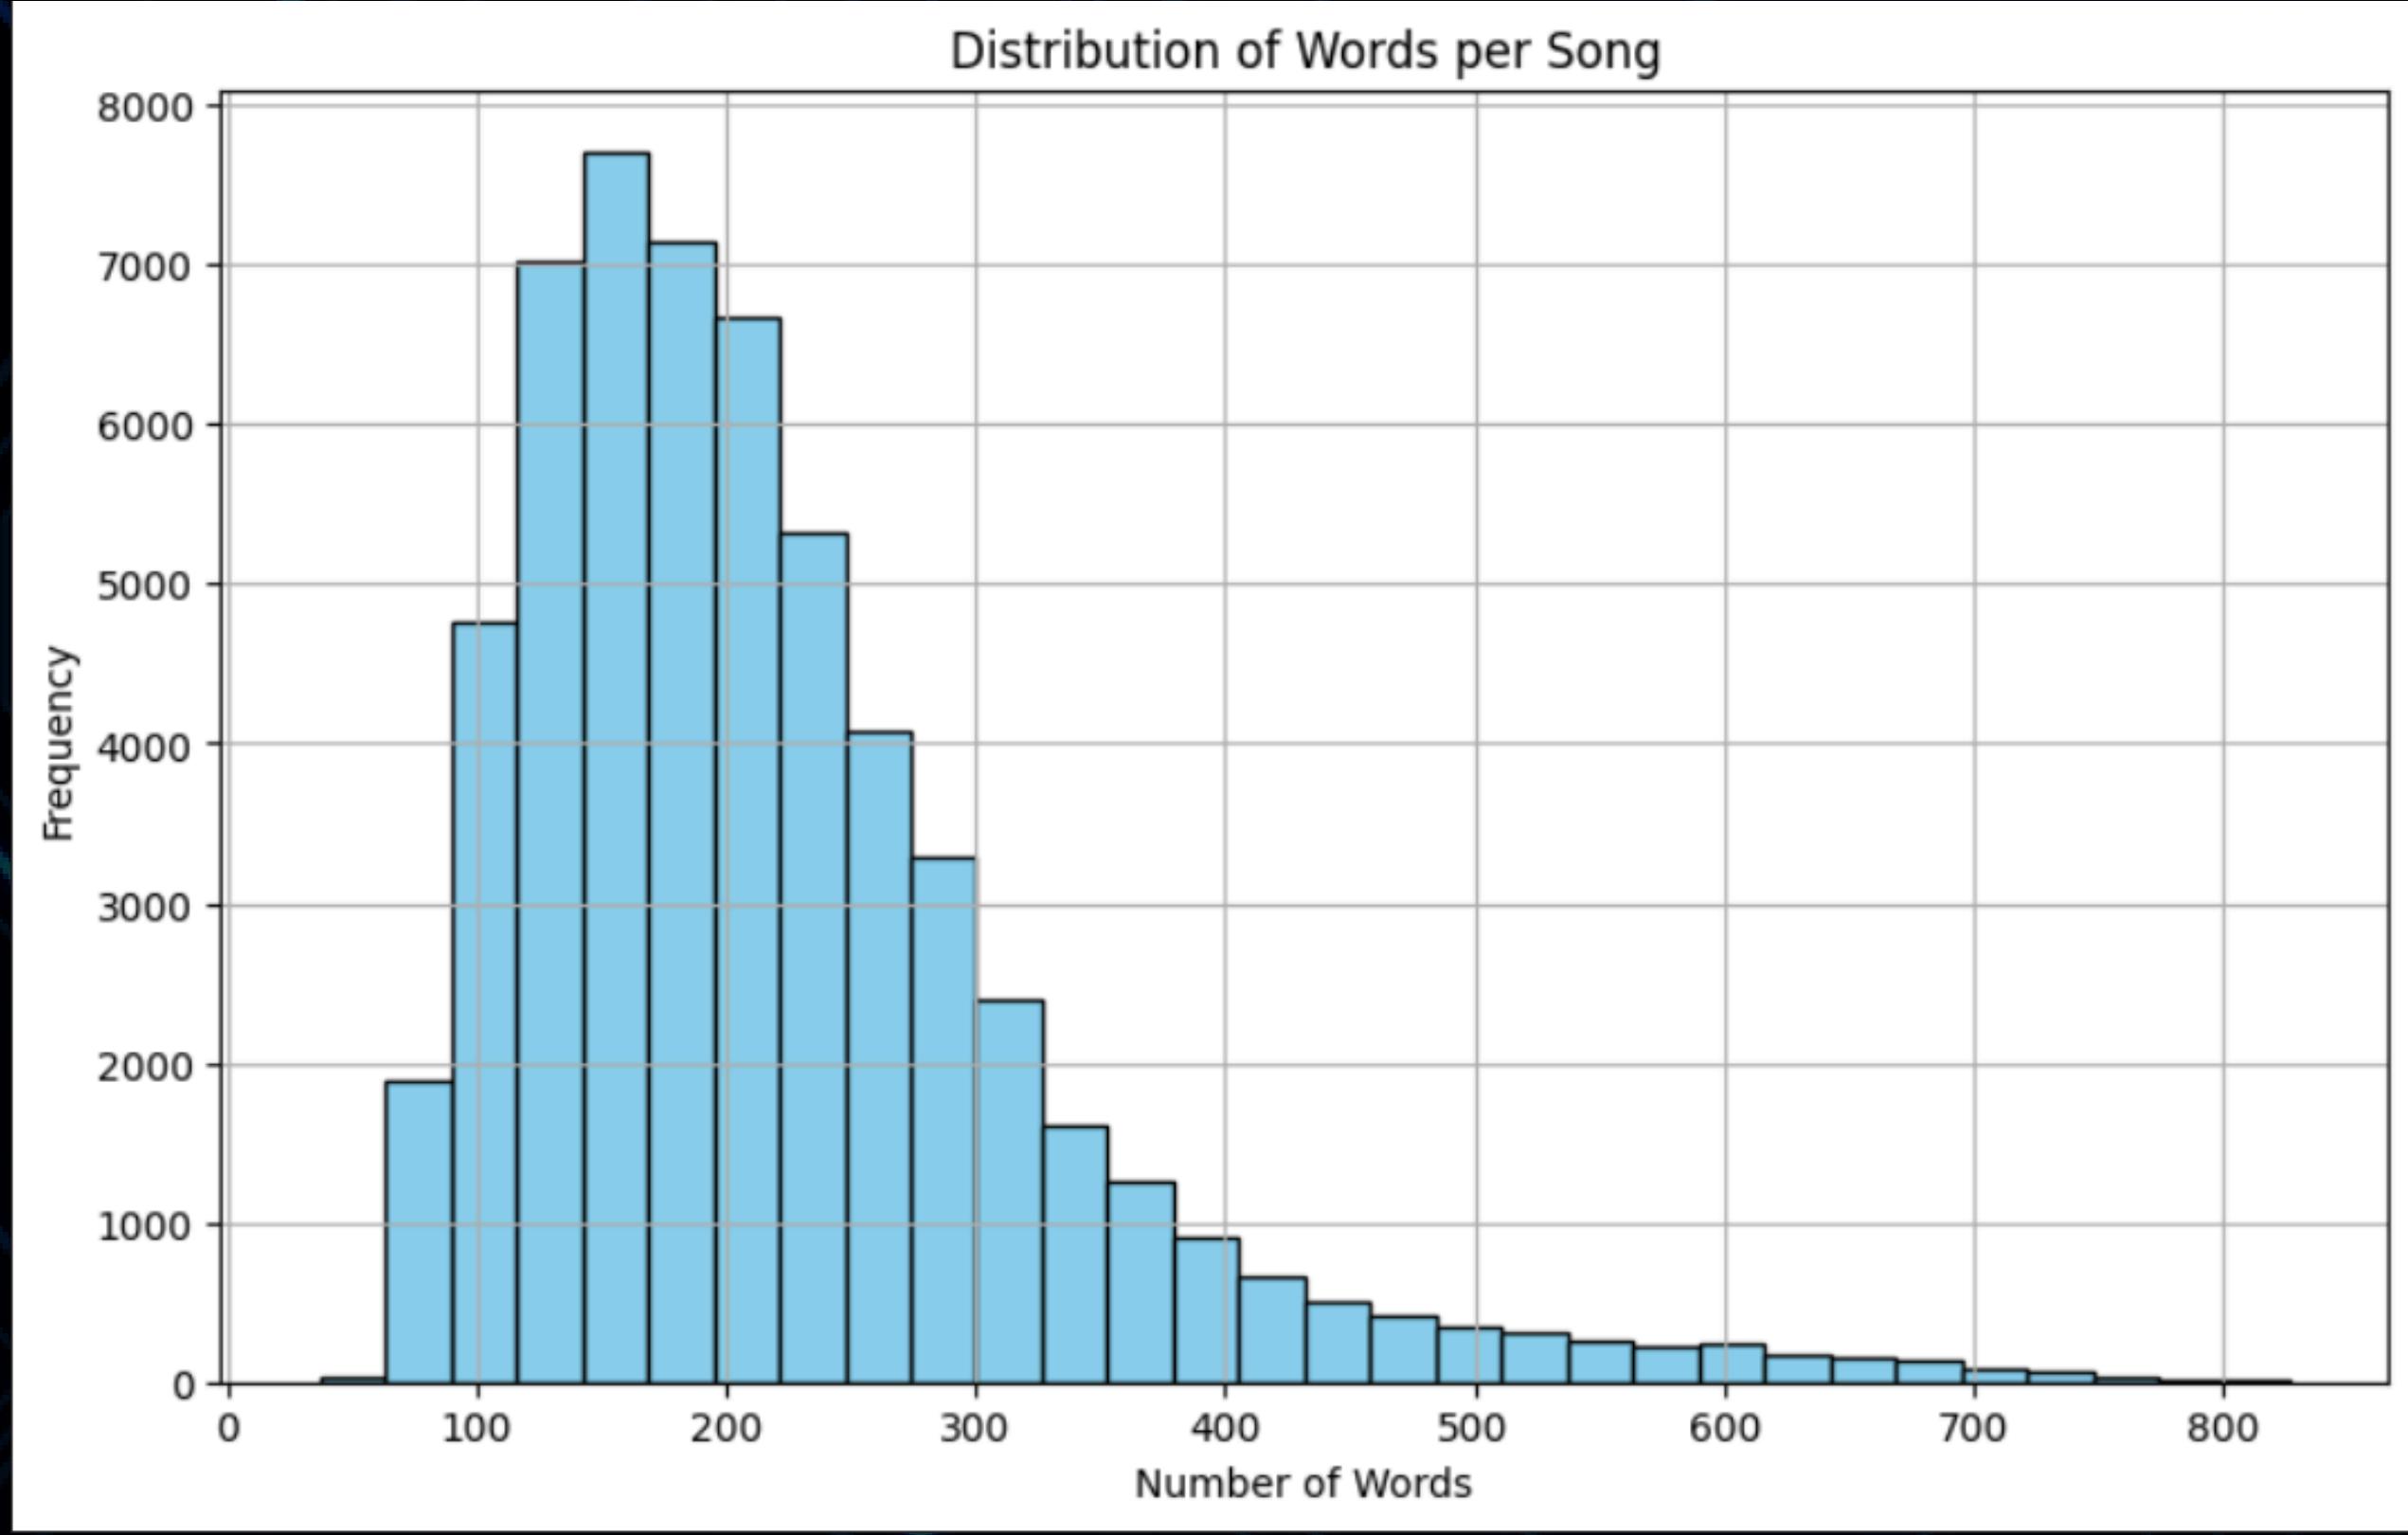
**FREE-FIREZONE**

# DATA PRE-PROCESSING

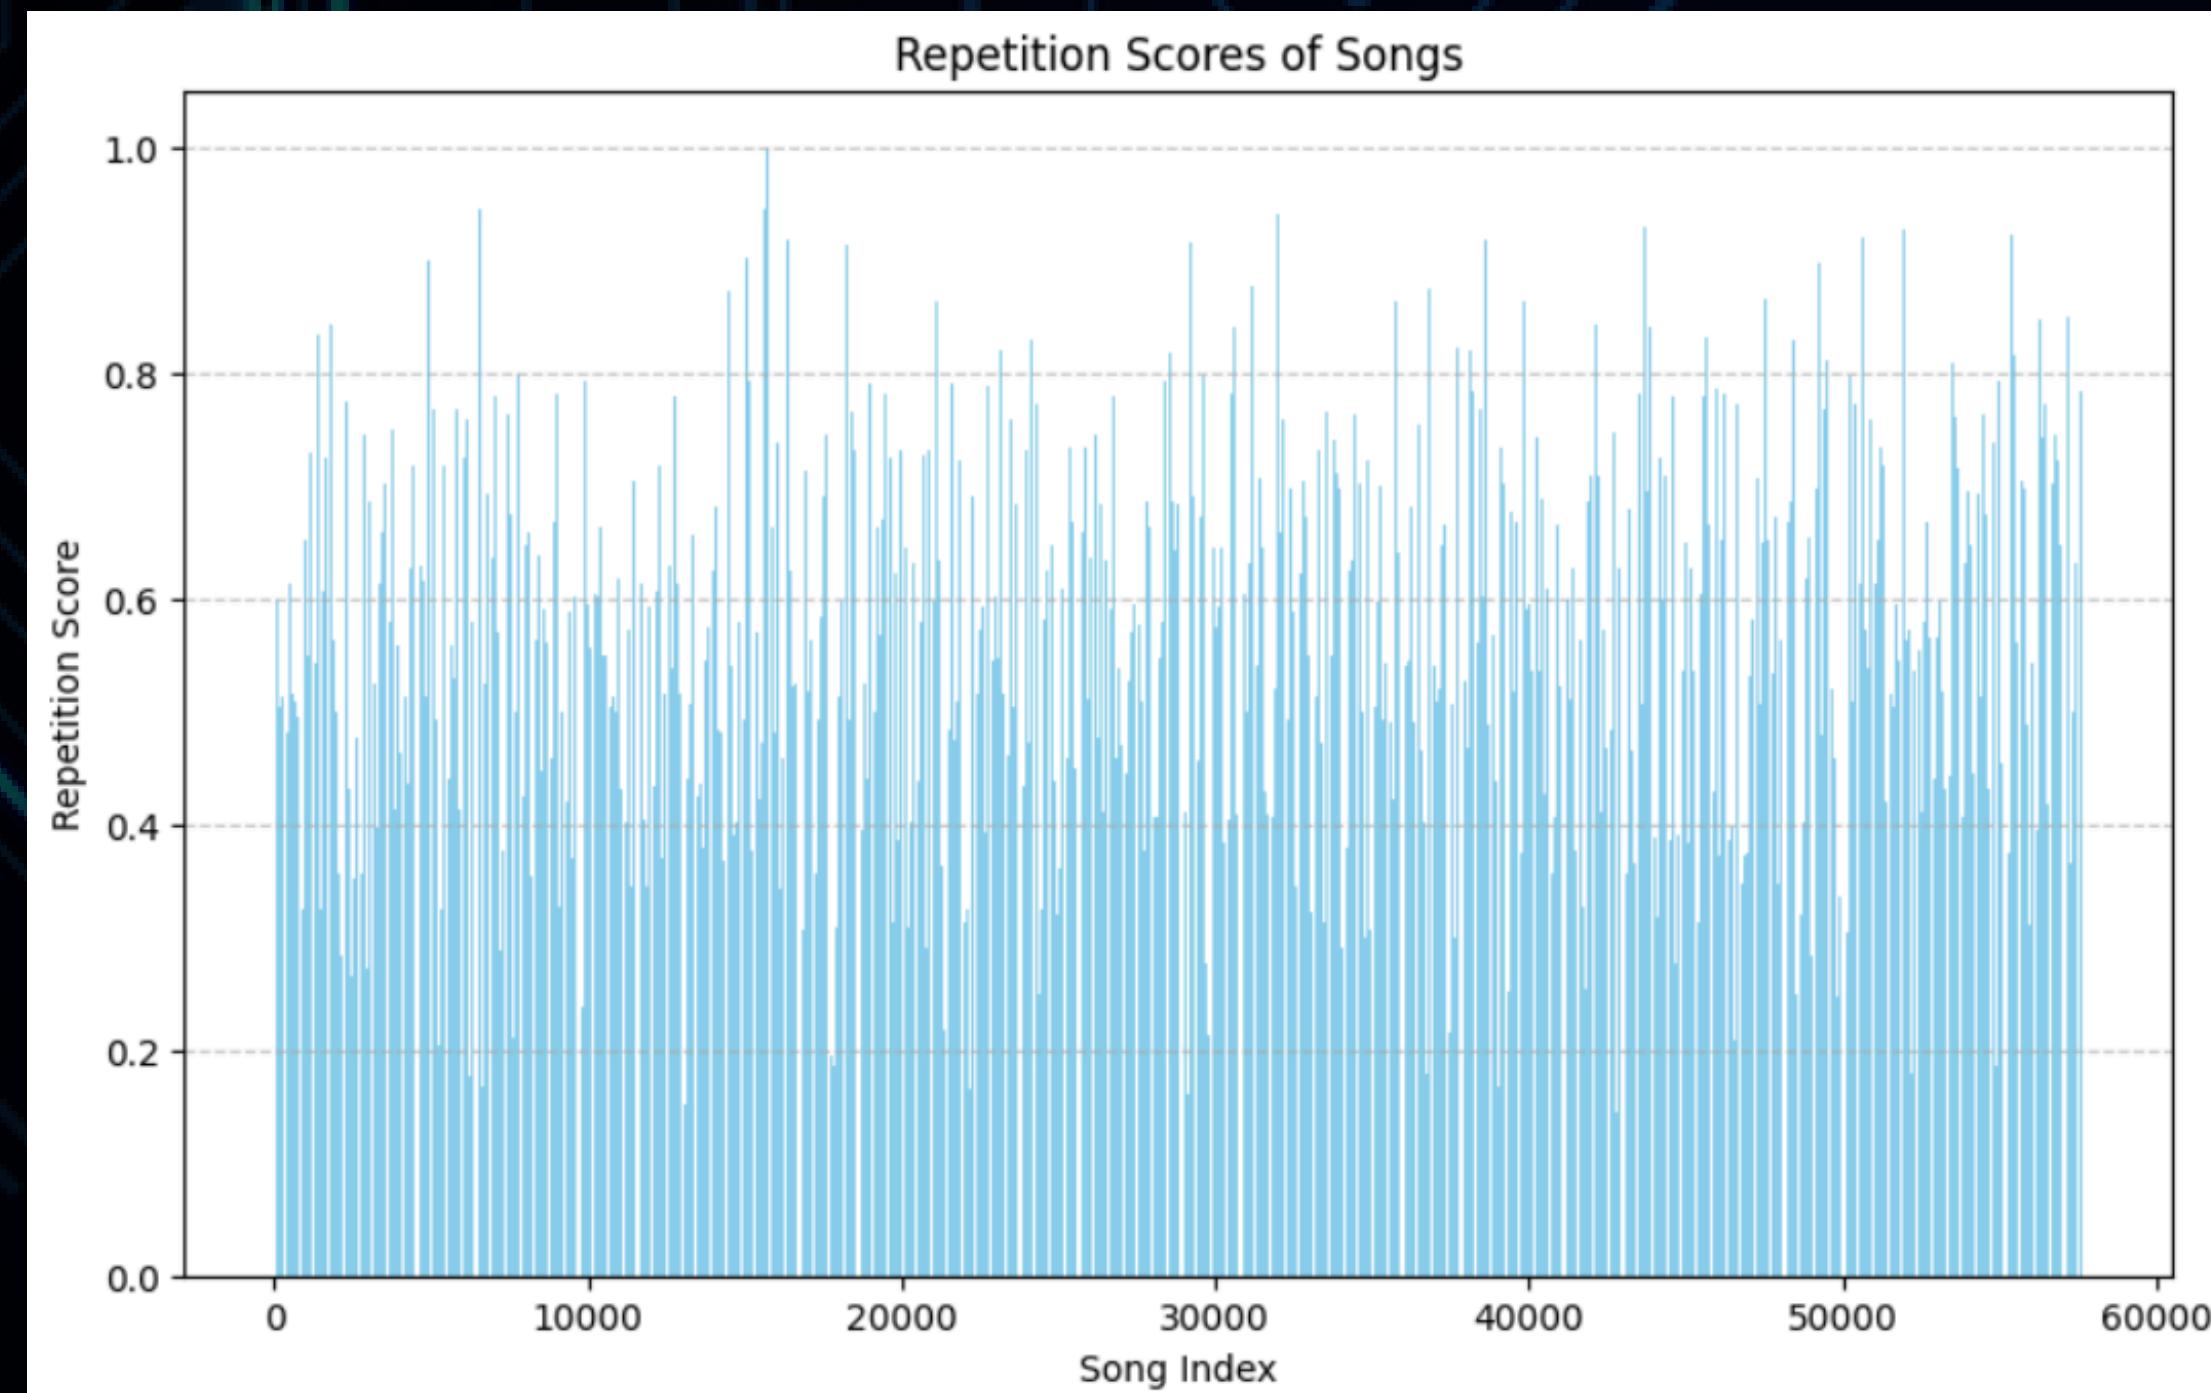
- Duplicates check
- Tokenization
- Stop word removal
- TF-IDF
- Word Cloud
- Some Graphs

## Top 30 Artists by Song Frequency

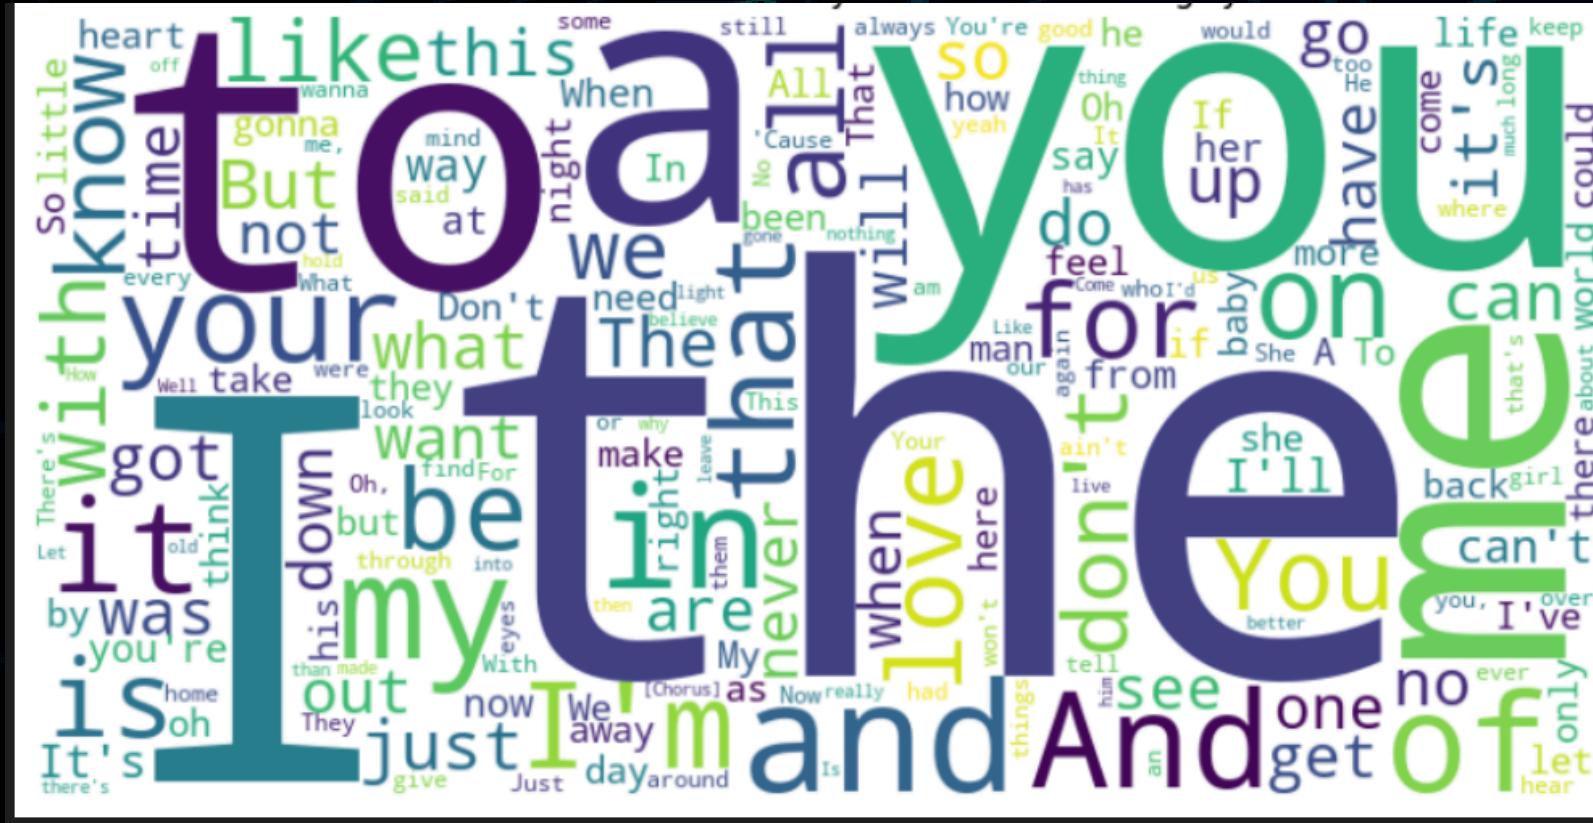




**Repetition Score** =  $\frac{\text{Number of Repeated Words}}{\text{Total Number of Words}}$



# BEFORE DATA CLEANING



# AFTER DATA CLEANING



AFFIRMATIVE

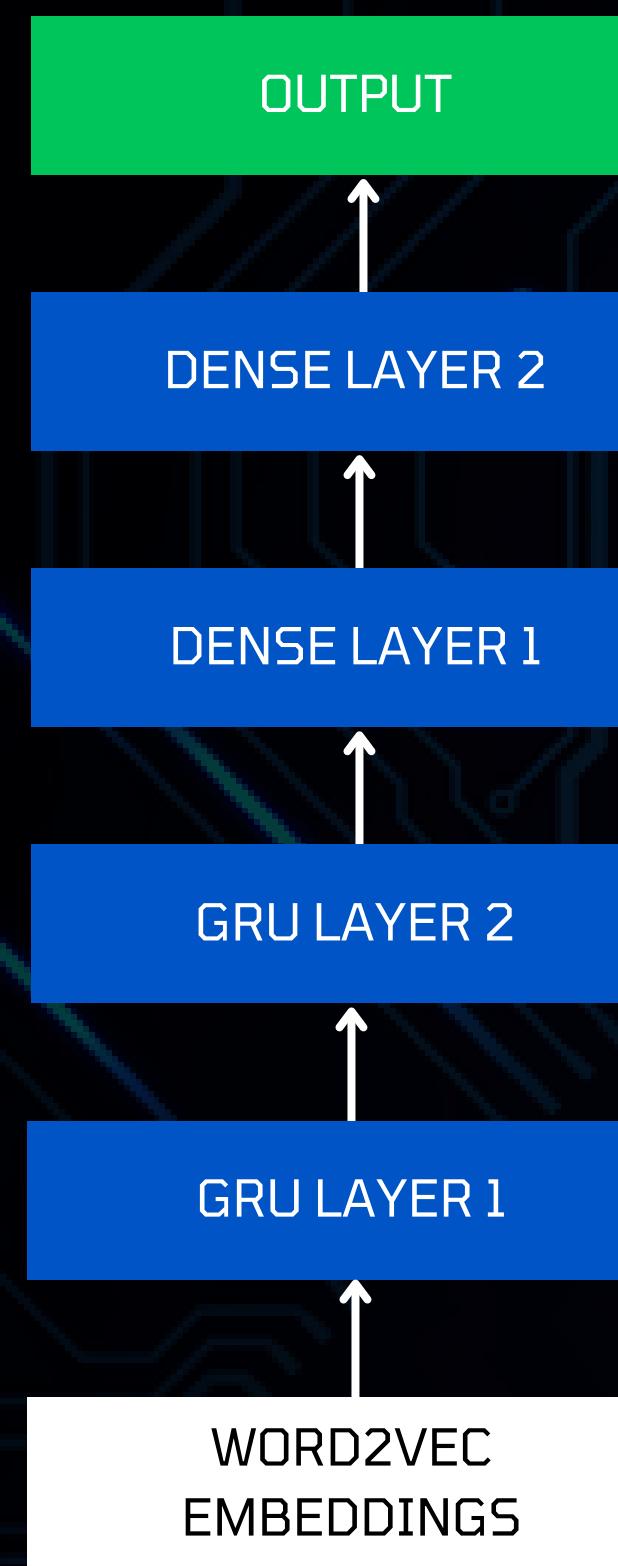
# APPROACHES

Our created  
model/pipeline

Pre-trained  
Model

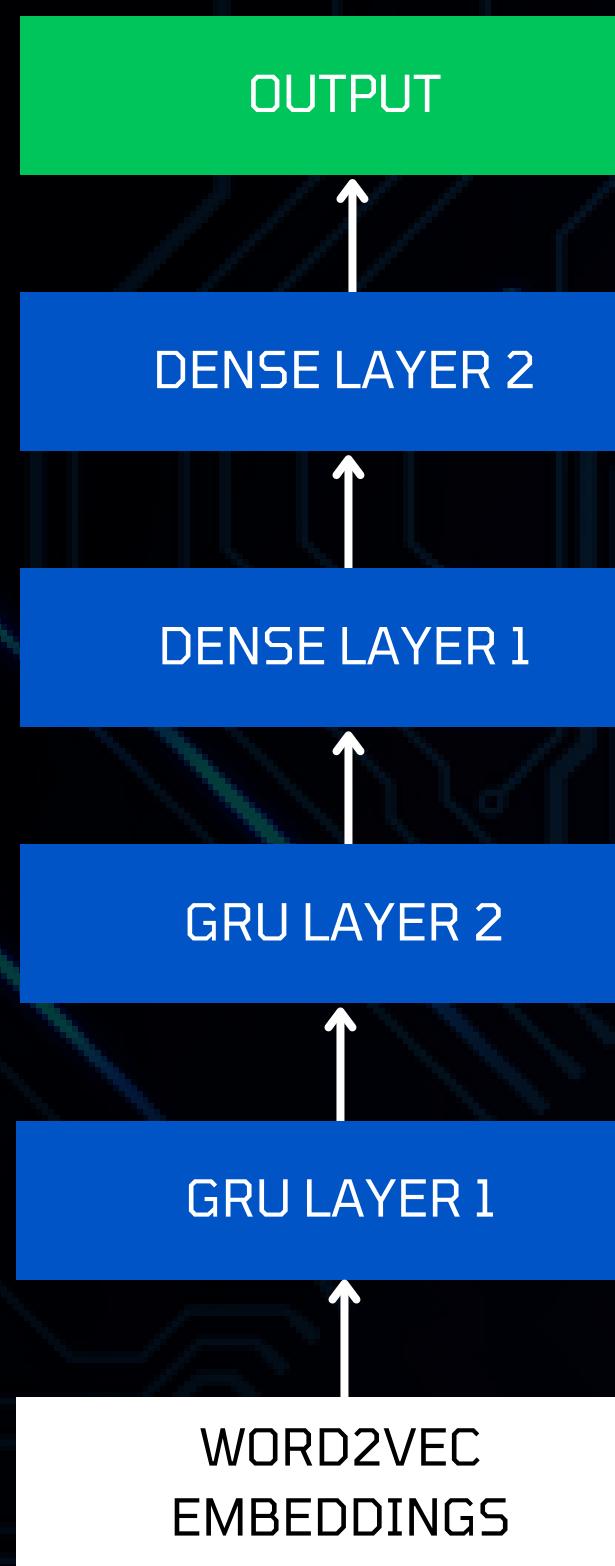


# CREATED MODEL ARCHITECTURE



- **Gated Recurrent Units (GRUs):**
  - Address vanishing gradient problem in RNNs.
  - Two GRU layers capture short-term & long-term dependencies:
    - First GRU (`return_sequences=True`) outputs a sequence.
    - Second GRU (`return_sequences=False`) summarizes the sequence.
- **Hidden Layer for Feature Extraction:**
  - Dense layer with 64 units & ReLU activation extracts features from GRU outputs.
- **Output Layer and Vocabulary:**
  - Final Dense layer with softmax activation predicts next word.
  - Number of units matches vocabulary size (all unique words).
  - Softmax activation gives probability distribution over vocabulary.

# CREATED MODEL ARCHITECTURE



Model: "sequential\_7"

Layer (type)	Output Shape	Param #
gru_14 (GRU)	(None, 224, 128)	53760
gru_15 (GRU)	(None, 128)	99072
dense_14 (Dense)	(None, 64)	8256
dense_15 (Dense)	(None, 109)	7085

Total params: 168173 (656.93 KB)

Trainable params: 168173 (656.93 KB)

Non-trainable params: 0 (0.00 Byte)

# RESULTS AND FINDINGS

Including 1 artist  
“Donna Summer”

```
Epoch 1/10
5/5 [=====] - 5s 123ms/step - loss: 0.6901 - accuracy: 0.0072
Epoch 2/10
5/5 [=====] - 1s 125ms/step - loss: 0.6670 - accuracy: 0.0072
Epoch 3/10
5/5 [=====] - 1s 122ms/step - loss: 0.5388 - accuracy: 0.0072
Epoch 4/10
5/5 [=====] - 1s 119ms/step - loss: 0.2464 - accuracy: 0.0072
Epoch 5/10
5/5 [=====] - 1s 119ms/step - loss: 0.1530 - accuracy: 0.0072
Epoch 6/10
5/5 [=====] - 1s 118ms/step - loss: 0.0973 - accuracy: 0.0072
Epoch 7/10
5/5 [=====] - 1s 120ms/step - loss: 0.0674 - accuracy: 0.0072
Epoch 8/10
5/5 [=====] - 1s 121ms/step - loss: 0.0559 - accuracy: 0.0290
Epoch 9/10
5/5 [=====] - 1s 127ms/step - loss: 0.0523 - accuracy: 0.1087
Epoch 10/10
5/5 [=====] - 1s 120ms/step - loss: 0.0509 - accuracy: 0.1087
```

# RESULTS AND FINDINGS

```
Epoch 1/10  
80/80 [=====] - 13s 110ms/step - loss: 0.1655 - accuracy: 0.0208  
Epoch 2/10  
80/80 [=====] - 9s 107ms/step - loss: 0.0073 - accuracy: 0.0513  
Epoch 3/10  
80/80 [=====] - 9s 109ms/step - loss: 0.0072 - accuracy: 0.0431  
Epoch 4/10  
80/80 [=====] - 9s 113ms/step - loss: 0.0072 - accuracy: 0.0466  
Epoch 5/10  
80/80 [=====] - 9s 110ms/step - loss: 0.0072 - accuracy: 0.0478  
Epoch 6/10  
80/80 [=====] - 9s 110ms/step - loss: 0.0072 - accuracy: 0.0462  
Epoch 7/10  
80/80 [=====] - 9s 112ms/step - loss: 0.0072 - accuracy: 0.0517  
Epoch 8/10  
80/80 [=====] - 9s 111ms/step - loss: 0.0072 - accuracy: 0.0493  
Epoch 9/10  
80/80 [=====] - 9s 113ms/step - loss: 0.0072 - accuracy: 0.0501  
Epoch 10/10  
80/80 [=====] - 9s 113ms/step - loss: 0.0072 - accuracy: 0.0525
```

Including more artists

# TESTING OUR MODEL

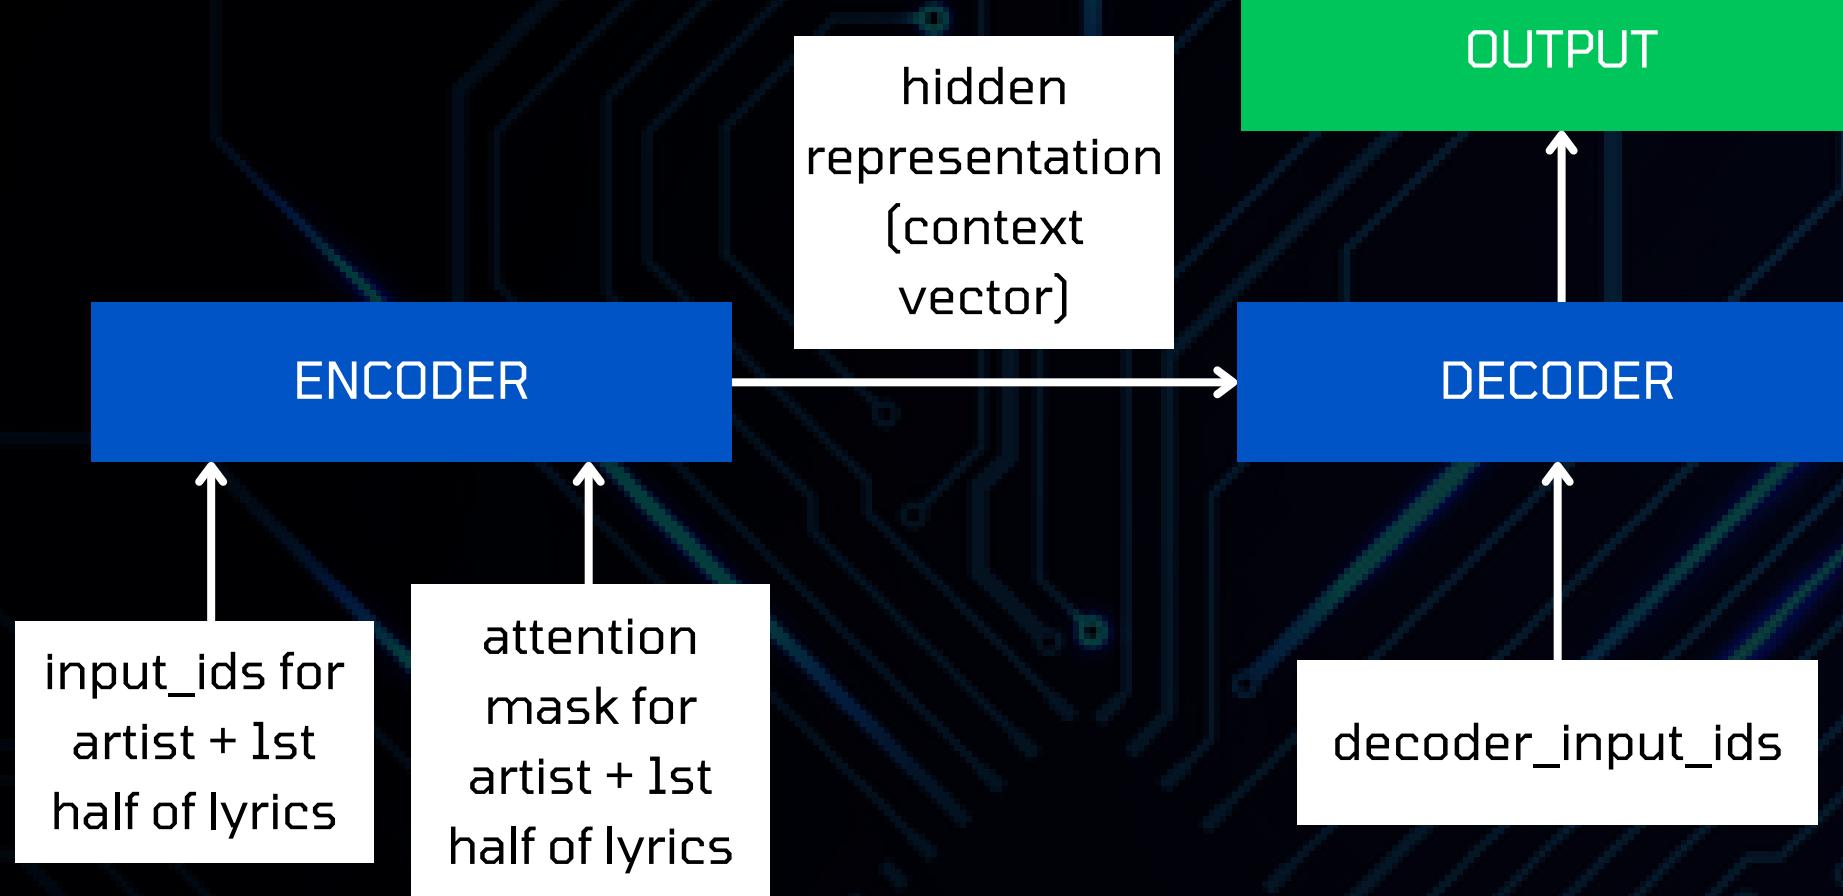
- Identical words due to the low accuracy
  - It's important to note that the quality of the generated text is directly proportional to the model's accuracy.

```
[['queen', 'day', 'queen', 'night', 'dressed', 'head', 'toe']]  
1/1 0s 33ms/step  
1/1 0s 32ms/step  
1/1 0s 36ms/step  
1/1 0s 28ms/step  
1/1 0s 63ms/step  
1/1 0s 29ms/step  
1/1 0s 29ms/step  
1/1 0s 31ms/step  
1/1 0s 30ms/step  
1/1 0s 37ms/step  
queen day queen night dressed head toe oh oh oh oh oh oh oh oh oh
```

# ADVANTAGES/DISADVANTAGES

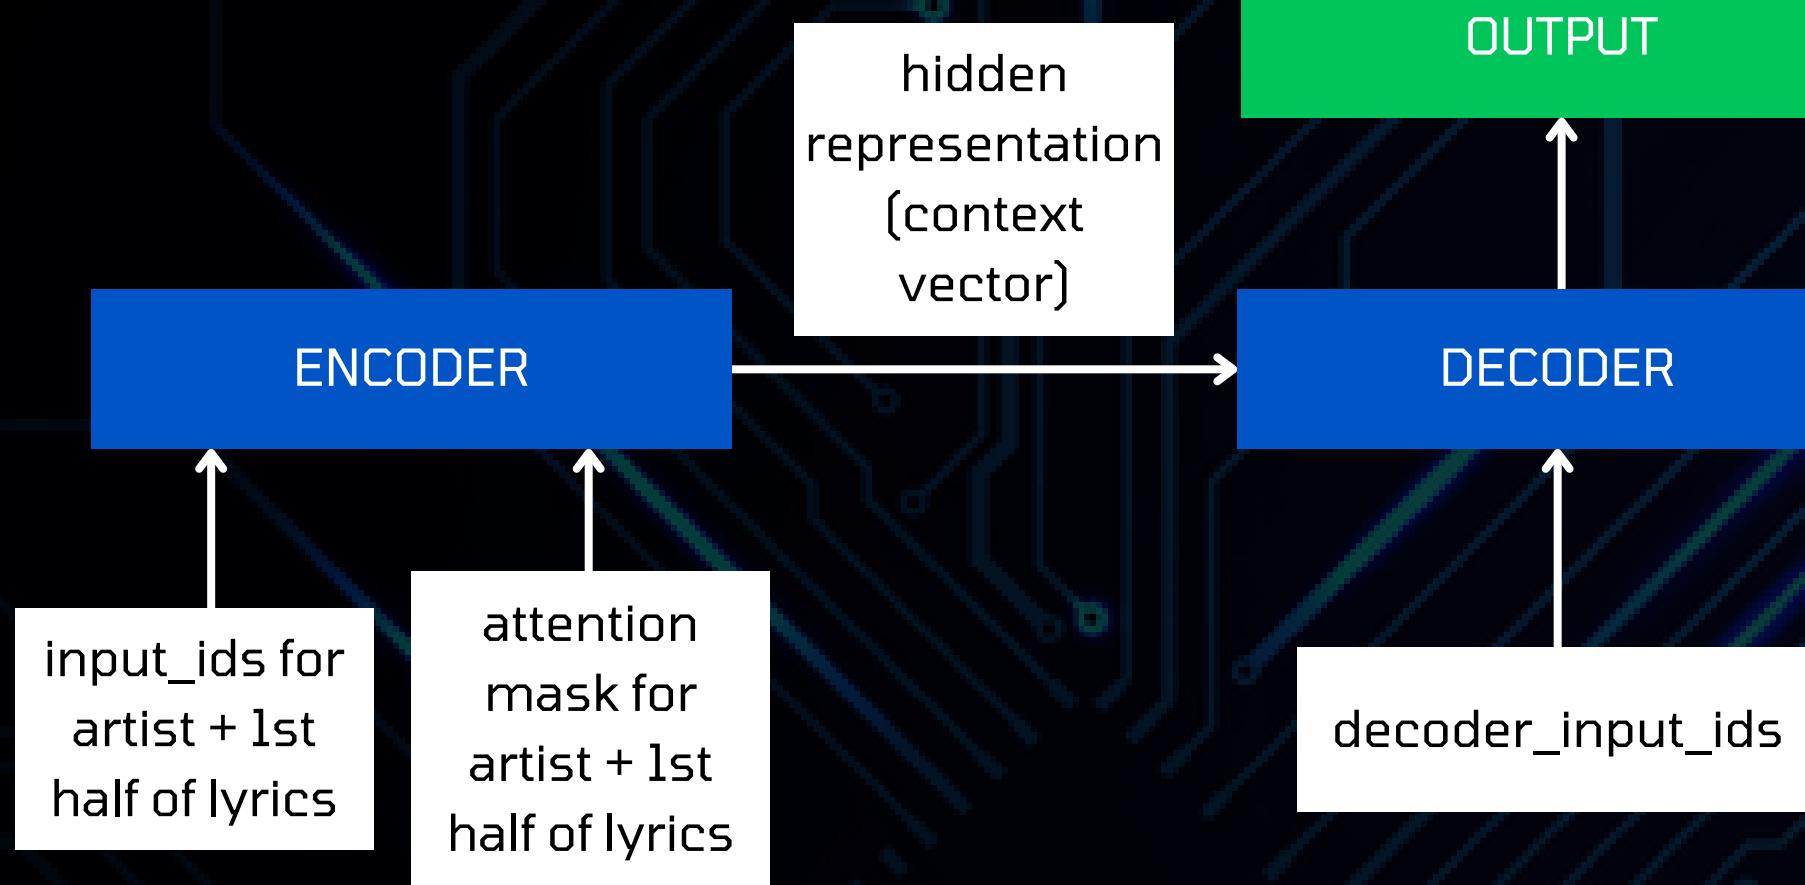
- *Advantages:*
  - Easier to understand and analyze how the model processes sequences due to its simpler structure.
  - Having the full control over the model architecture and can experiment with hyperparameters for optimization.
  - Training a custom GRU model might require less computational power compared to pre-trained Transformers.
- *Disadvantages and Limitations:*
  - May not achieve the same level of performance as pre-trained Transformers in complex tasks like lyric generation.

# T5 PRE-TRAINED MODEL



- **T5:**
  - Handles various text tasks, including lyric generation.
  - Takes additional instructions to tailor outputs.
- **Conditional T5:**
  - Provide instructions alongside the main input sequence.
- **Encoder-Decoder Architecture:**
  - Encoder: Processes input & conditions, captures meaning.
  - Decoder: Generates target text (lyrics) influenced by conditions.
- **Fine-tuning on Lyric Data:**
  - Train a pre-trained T5 model on lyrics with conditions.
  - Improves understanding of lyric patterns and vocabulary.

# T5 PRE-TRAINED MODEL



Model: "tft5\_for\_conditional\_generation"

Layer (type)	Output Shape	Param #
shared (Embedding)	multiple	16449536
encoder (TFT5MainLayer)	multiple	35330816
decoder (TFT5MainLayer)	multiple	41625344
<hr/>		
Total params: 60506624 (230.81 MB)		
Trainable params: 60506624 (230.81 MB)		
Non-trainable params: 0 (0.00 Byte)		

# TS, PRE-TRAINED MODEL

- Approach A:
  - **model.fit()**
    - Converted data to NumPy format for efficiency in machine learning.
    - Defined optimizer, loss function, and accuracy metrics.
    - Compiled the model with defined elements for training.
    - Trained on provided data and monitored performance on unseen test data.

# TF2 PRE-TRAINED MODEL

- Approach B:
  - **training a sequence-to-sequence model with Hugging Face Transformers (for loop)**
    - Converted dataset to TensorFlow format for efficient processing and shuffling.
    - Defined train step: *process batch data, calculate loss, update gradients, update weights, update accuracy.*
    - Defined evaluate step: *process batch data, calculate loss, update accuracy (for unseen data).*
    - Trained in epochs: *iteratively train on batches, evaluate on unseen data after each epoch.*

# RESULTS AND FINDINGS (APPROACH 1)

On Top 5 artists

```
[68]: model.fit(train_dataset, validation_data=test_dataset,batch_size=16, epochs=3)

Epoch 1/3
WARNING: AutoGraph could not transform <function infer_framework at 0x7dae71a17400> and will run it as-is.
Cause: for/else statement not yet supported
To silence this warning, decorate the function with @tf.autograph.experimental.do_not_convert
48/48 [=====] - 164s 1s/step - loss: 4.4051 - sparse_categorical_accuracy: 0.4947 - val_loss: 1.2440 - val_sparse_categorical_accuracy: 0.7962
Epoch 2/3
48/48 [=====] - 35s 728ms/step - loss: 1.6087 - sparse_categorical_accuracy: 0.7422 - val_loss: 1.1684 - val_sparse_categorical_accuracy: 0.8048
Epoch 3/3
48/48 [=====] - 35s 729ms/step - loss: 1.2830 - sparse_categorical_accuracy: 0.7863 - val_loss: 1.1071 - val_sparse_categorical_accuracy: 0.8108
[68]: <tf_keras.src.callbacks.History at 0x7dae2410c820>
```

# RESULTS AND FINDINGS (APPROACH 1)

Full Dataset on kaggle with GPU T4X2

The screenshot shows a Jupyter Notebook interface running on a Kaggle draft session. The session has been running for 6 hours and 39 minutes. The notebook contains two cells:

- [31]:  
model.compile(optimizer=optimizer, loss=loss, metrics=metrics)
- [32]:  
model.fit(train\_dataset, validation\_data=test\_dataset, batch\_size=16, epochs=3)  
The output of this cell shows training logs for three epochs. The accuracy values for epoch 3 are highlighted with a red box:
  - Epoch 1/3
  - Epoch 2/3
  - Epoch 3/3  
2883/2883 [=====] - 3798s 1s/step - loss: 1.1343 - sparse\_categorical\_accuracy: 0.8102 - val\_loss: 0.9327 - val\_sparse\_categorical\_accuracy: 0.8361
  - Epoch 3/3  
2883/2883 [=====] - 3772s 1s/step - loss: 0.9803 - sparse\_categorical\_accuracy: 0.8285 - val\_loss: 0.8984 - val\_sparse\_categorical\_accuracy: 0.8403
  - Epoch 3/3  
2883/2883 [=====] - 3774s 1s/step - loss: 0.9502 - sparse\_categorical\_accuracy: 0.8325 - val\_loss: 0.8818 - val\_sparse\_categorical\_accuracy: 0.8427

The session options panel on the right indicates the use of "GPU T4 x2".

# RESULTS AND FINDINGS (APPROACH 1)

Full Dataset on kaggle with GPU P100

```
[33]: model.fit(train_dataset, validation_data=test_dataset,batch_size=16, epochs=3)

Epoch 1/3
WARNING: AutoGraph could not transform <function infer_framework at 0x7b0f774a2dd0> and will run it as-is.
Cause: for/else statement not yet supported
To silence this warning, decorate the function with @tf.autograph.experimental.do_not_convert
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
I0000 00:00:1716047933.435422      89 device_compiler.h:186] Compiled cluster using XLA! This line is logged at most once for
the lifetime of the process.
2883/2883 [=====] - 2310s 764ms/step - loss: 1.1300 - sparse_categorical_accuracy: 0.8108 - val_loss:
0.9427 - val_sparse_categorical_accuracy: 0.8343
Epoch 2/3
2883/2883 [=====] - 2183s 757ms/step - loss: 0.9780 - sparse_categorical_accuracy: 0.8289 - val_loss:
0.9069 - val_sparse_categorical_accuracy: 0.8387
Epoch 3/3
2883/2883 [=====] - 2186s 758ms/step - loss: 0.9477 - sparse_categorical_accuracy: 0.8329 - val_loss:
0.8882 - val_sparse_categorical_accuracy: 0.8412
```

[33]: <tf\_keras.src.callbacks.History at 0x7b0ea44c1fc0>

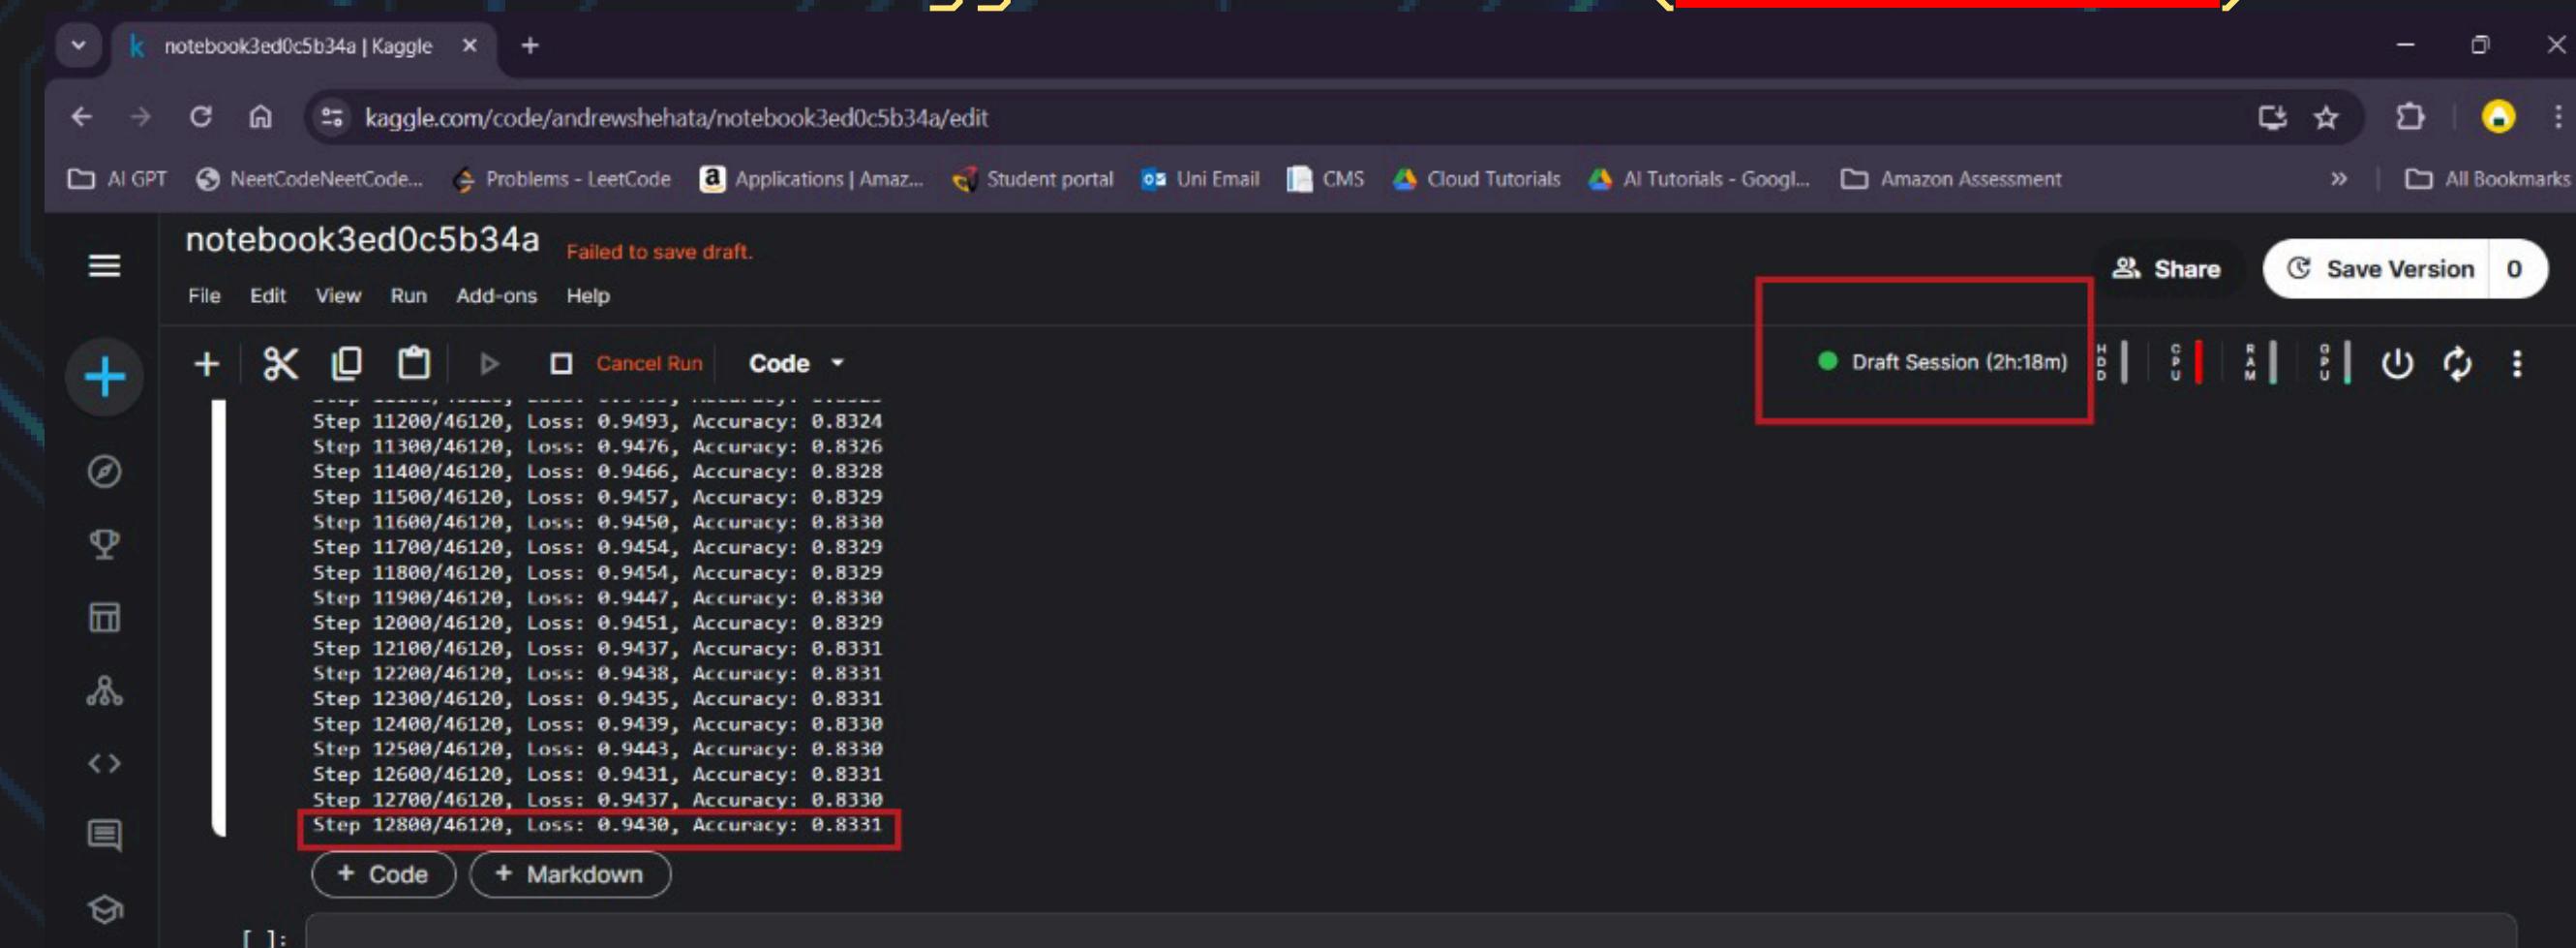
# RESULTS AND FINDINGS (APPROACH 2)

On Top 5 artists

```
TRAINING
Epoch 1/3
Step 100/754, Loss: 2.9951, Accuracy: 0.6227
Step 200/754, Loss: 2.0151, Accuracy: 0.7241
Step 300/754, Loss: 1.6826, Accuracy: 0.7564
Step 400/754, Loss: 1.5234, Accuracy: 0.7707
Step 500/754, Loss: 1.4276, Accuracy: 0.7794
Step 600/754, Loss: 1.3637, Accuracy: 0.7860
Step 700/754, Loss: 1.3252, Accuracy: 0.7892
EVALUATION
Validation Loss: 0.9789, Validation Accuracy: 0.8265
TRAINING
Epoch 2/3
Step 100/754, Loss: 0.9598, Accuracy: 0.8302
Step 200/754, Loss: 0.9571, Accuracy: 0.8305
Step 300/754, Loss: 0.9486, Accuracy: 0.8316
Step 400/754, Loss: 0.9558, Accuracy: 0.8298
Step 500/754, Loss: 0.9559, Accuracy: 0.8300
Step 600/754, Loss: 0.9609, Accuracy: 0.8292
Step 700/754, Loss: 0.9592, Accuracy: 0.8291
EVALUATION
Validation Loss: 0.9464, Validation Accuracy: 0.8313
TRAINING
Epoch 3/3
Step 100/754, Loss: 0.9786, Accuracy: 0.8276
...
Step 600/754, Loss: 0.9045, Accuracy: 0.8372
Step 700/754, Loss: 0.9134, Accuracy: 0.8352
EVALUATION
Validation Loss: 0.9322, Validation Accuracy: 0.8326
```

# RESULTS AND FINDINGS [APPROACH E]

Full Dataset on kaggle with GPU P100 (unsuccessful)

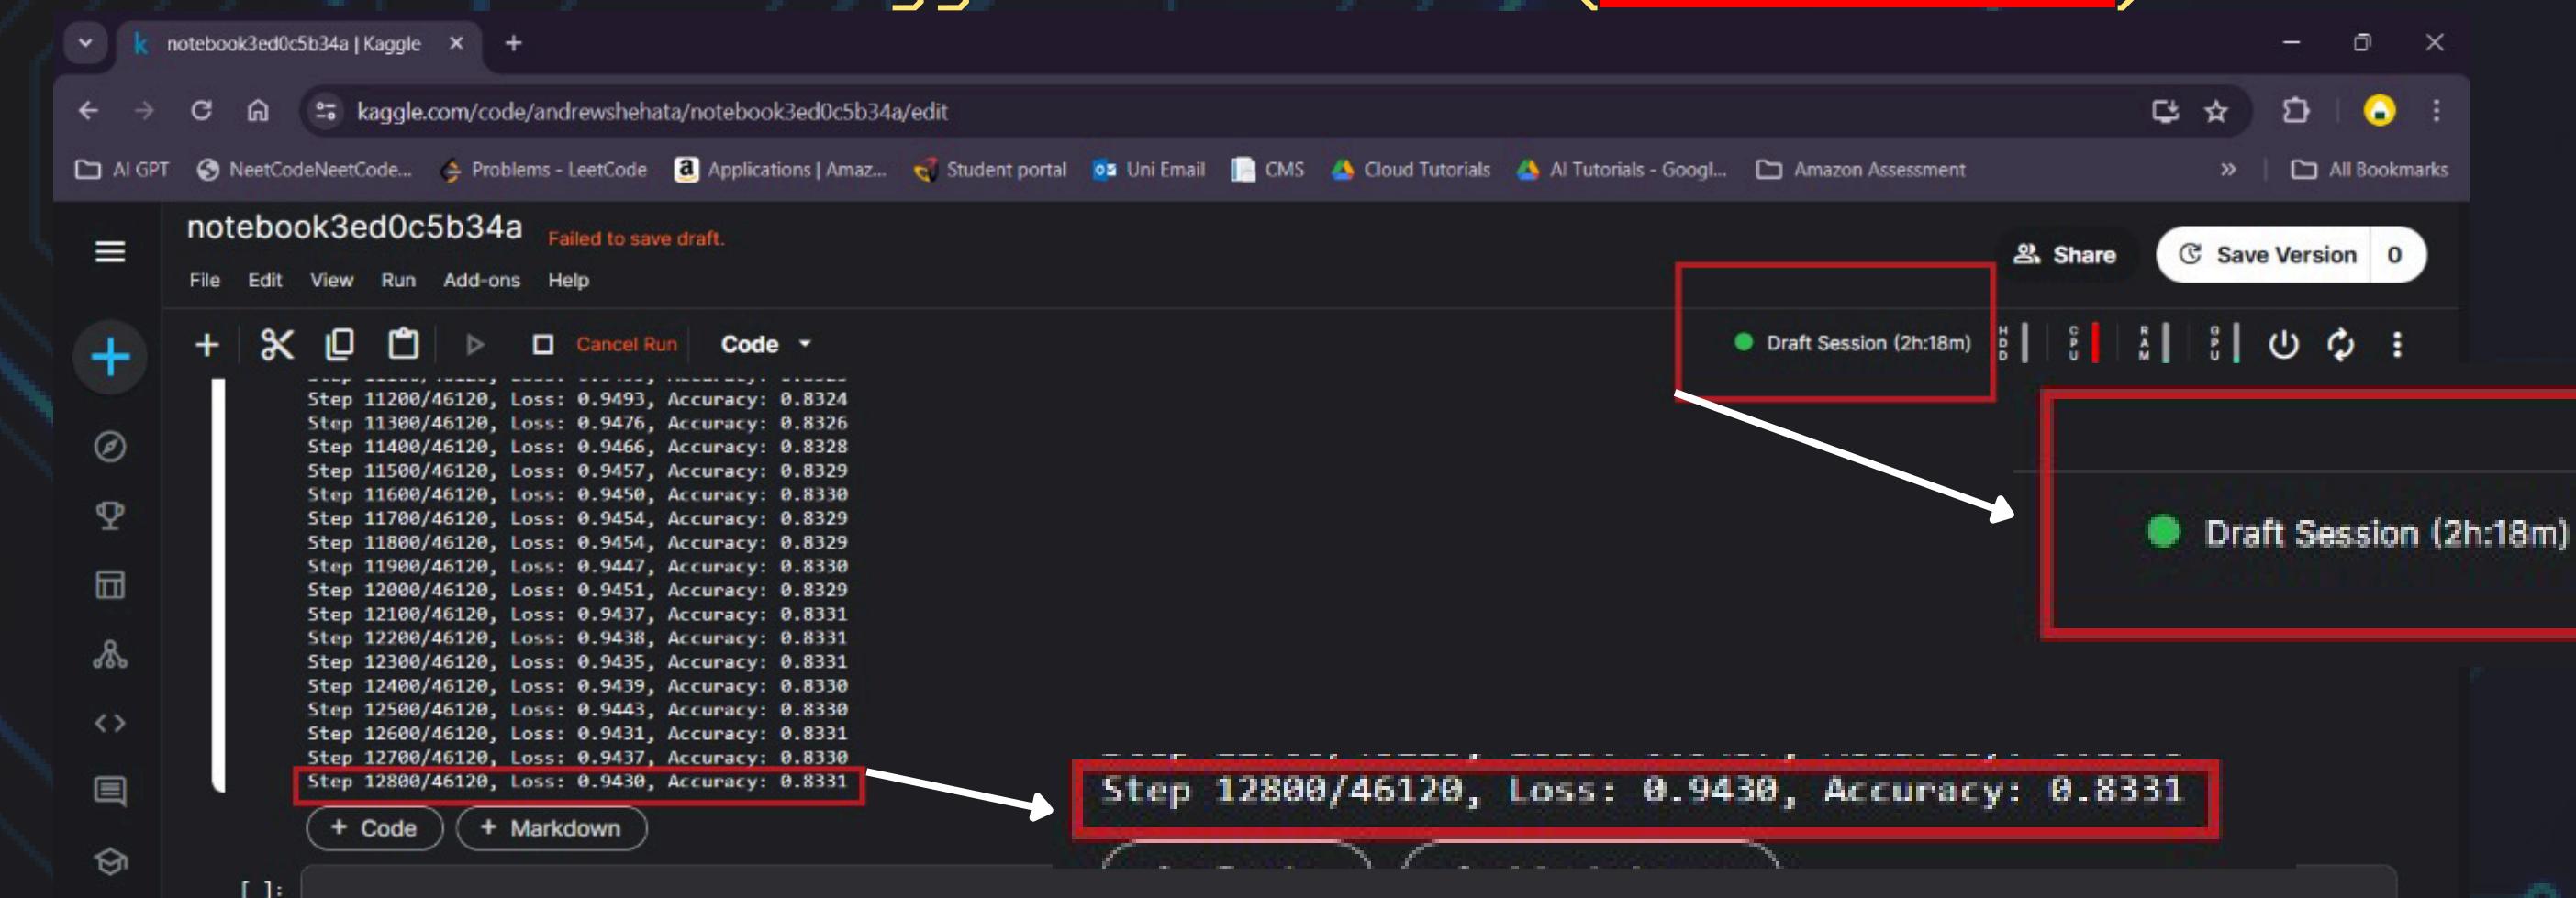


The screenshot shows a Jupyter Notebook interface on the Kaggle platform. The notebook title is "notebook3ed0c5b34a". The status bar at the top right indicates a "Draft Session (2h:18m)". The main code cell displays a series of training log entries:

```
Step 11200/46120, Loss: 0.9493, Accuracy: 0.8324
Step 11300/46120, Loss: 0.9476, Accuracy: 0.8326
Step 11400/46120, Loss: 0.9466, Accuracy: 0.8328
Step 11500/46120, Loss: 0.9457, Accuracy: 0.8329
Step 11600/46120, Loss: 0.9458, Accuracy: 0.8330
Step 11700/46120, Loss: 0.9454, Accuracy: 0.8329
Step 11800/46120, Loss: 0.9454, Accuracy: 0.8329
Step 11900/46120, Loss: 0.9447, Accuracy: 0.8330
Step 12000/46120, Loss: 0.9451, Accuracy: 0.8329
Step 12100/46120, Loss: 0.9437, Accuracy: 0.8331
Step 12200/46120, Loss: 0.9438, Accuracy: 0.8331
Step 12300/46120, Loss: 0.9435, Accuracy: 0.8331
Step 12400/46120, Loss: 0.9439, Accuracy: 0.8330
Step 12500/46120, Loss: 0.9443, Accuracy: 0.8330
Step 12600/46120, Loss: 0.9431, Accuracy: 0.8331
Step 12700/46120, Loss: 0.9437, Accuracy: 0.8330
Step 12800/46120, Loss: 0.9430, Accuracy: 0.8331
```

# RESULTS AND FINDINGS [APPROACH E]

Full Dataset on kaggle with GPU P100 (unsuccessful)



The screenshot shows a Jupyter Notebook interface on the Kaggle platform. The notebook title is "notebook3ed0c5b34a". The code cell contains a series of training log entries:

```
Step 11200/46120, Loss: 0.9493, Accuracy: 0.8324
Step 11300/46120, Loss: 0.9476, Accuracy: 0.8326
Step 11400/46120, Loss: 0.9466, Accuracy: 0.8328
Step 11500/46120, Loss: 0.9457, Accuracy: 0.8329
Step 11600/46120, Loss: 0.9458, Accuracy: 0.8330
Step 11700/46120, Loss: 0.9454, Accuracy: 0.8329
Step 11800/46120, Loss: 0.9454, Accuracy: 0.8329
Step 11900/46120, Loss: 0.9447, Accuracy: 0.8330
Step 12000/46120, Loss: 0.9451, Accuracy: 0.8329
Step 12100/46120, Loss: 0.9437, Accuracy: 0.8331
Step 12200/46120, Loss: 0.9438, Accuracy: 0.8331
Step 12300/46120, Loss: 0.9435, Accuracy: 0.8331
Step 12400/46120, Loss: 0.9439, Accuracy: 0.8330
Step 12500/46120, Loss: 0.9443, Accuracy: 0.8330
Step 12600/46120, Loss: 0.9431, Accuracy: 0.8331
Step 12700/46120, Loss: 0.9437, Accuracy: 0.8330
Step 12800/46120, Loss: 0.9430, Accuracy: 0.8331
```

The session status bar indicates "Draft Session (2h:18m)". A red box highlights the final log entry: "Step 12800/46120, Loss: 0.9430, Accuracy: 0.8331". A white arrow points from this box to a larger red box containing the same text. Another white arrow points from the bottom right of the log area to the session status bar.

# TESTING OUR MODEL

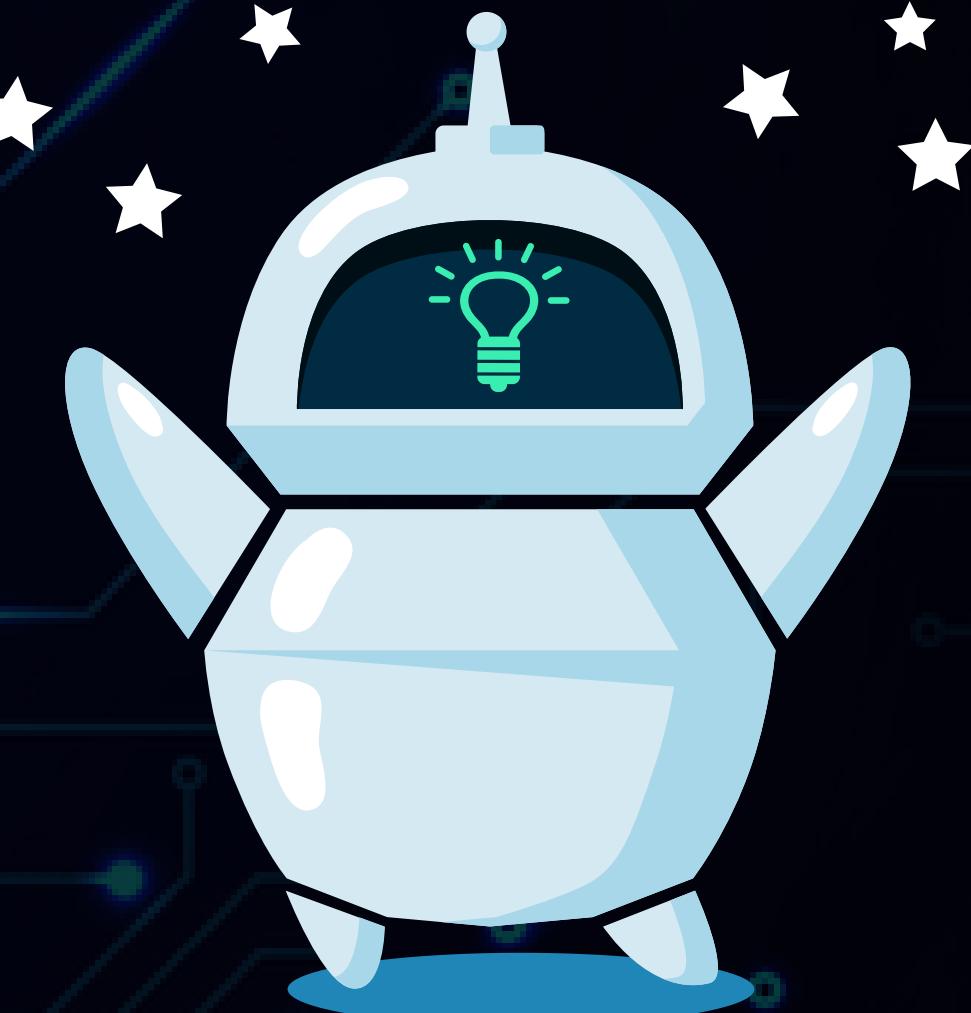
- Feed lyric snippet to tokenizer (text to numbers).
- Use model to generate continuation with attributes:
  - *max new tokens*: limit generated text length (e.g., 40 words).
  - *do sample*: inject randomness for diverse outputs.
  - *top k*: consider only top 30 probable tokens for next word.
  - *top p*: balance high-probability with some exploration.
- Decode generated sequence back to human-readable text (numbers to text), excluding special tokens.

# ADVANTAGES/DISADVANTAGES

- *Advantages:*
  - Leverages the pre-trained knowledge of TFT5, potentially leading to better lyric generation quality.
- *Disadvantages and Limitations:*
  - Understanding how a pre-trained Transformer arrives at its predictions can be challenging.
  - Have less control over the overall model architecture as you're primarily fine-tuning an existing model.
  - Training even with fine-tuning might require more computational resources compared to a custom GRU model.

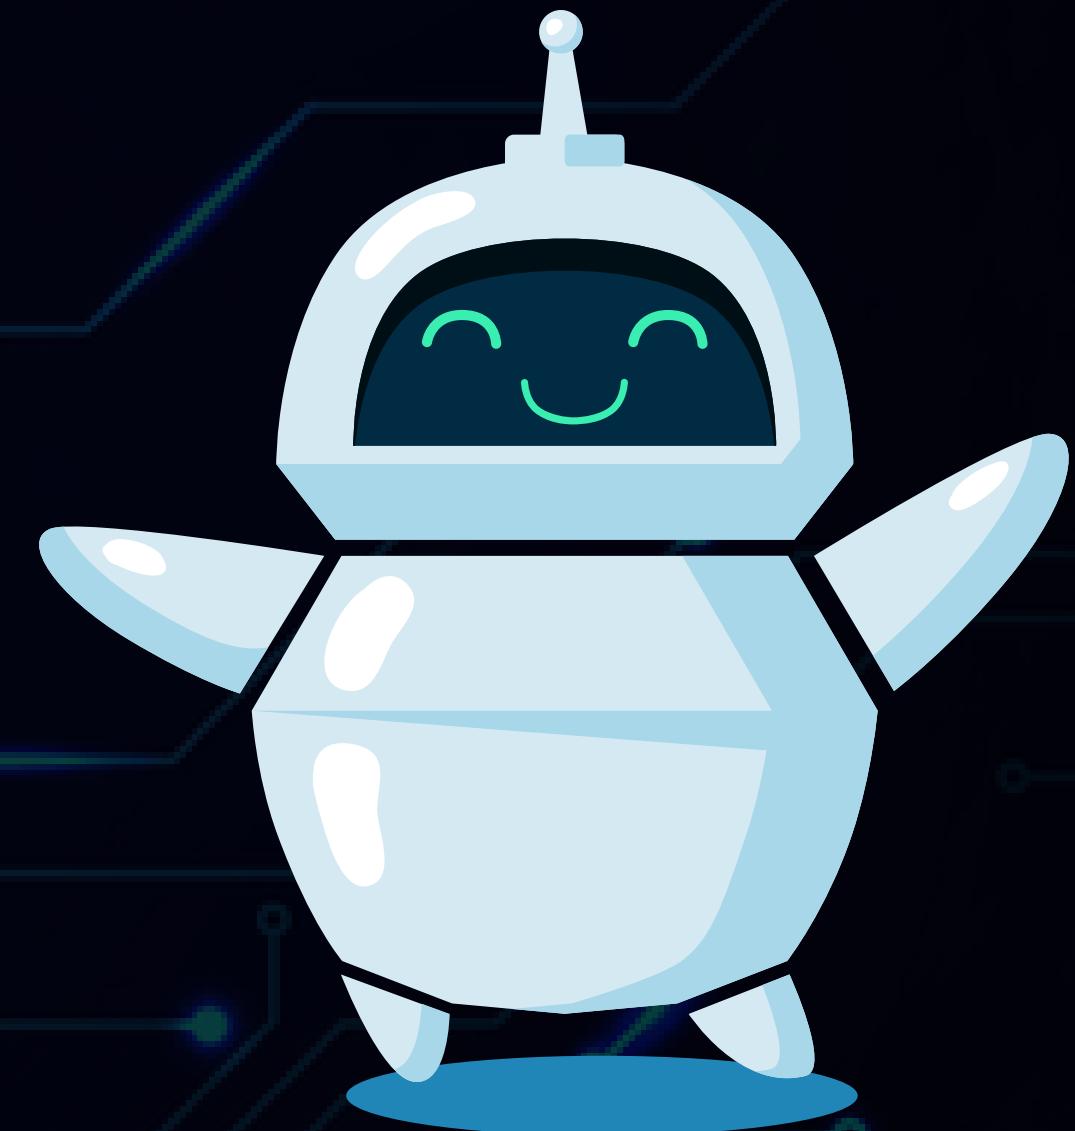
# CHALLENGES AND LIMITATIONS

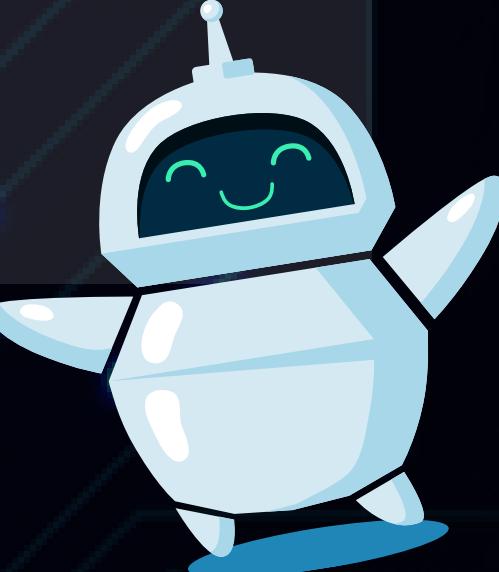
- **Small Dataset (56,000 entries):**
  - Training takes too long, especially with small batches for better understanding.
  - Lowers accuracy when not using a pre-trained model.
- **Hardware Constraints (laptops):**
  - Insufficient RAM and CPU lead to crashes during training .
- **Spelling Check Limitation:**
  - Takes too long (4 hours) for the entire dataset.
  - Song lyrics may not require strict spelling (e.g., "oh" vs "ohhhhh").
- **Sparse Feature Set:**
  - Lacks information like genre, duration, etc.
  - Needs web scraping or external data integration for better analysis.



# CONCLUSION

- **Interpretability & Customization:** Choose a custom GRU model (but expect lower performance and longer training).
- **Speed & Performance:** Use `model.fit` with a pre-trained TFT5 (balances ease of use with results).
- **Maximum Control & Expertise:** Implement a custom training loop with TFT5 (requires expertise but offers most control).





THANK YOU  
ANY QUESTIONS?